

Neil Kale

nkale@cs.cmu.edu | (508) 936-3455 | neilkale.github.io | www.linkedin.com/in/neil-kale/

Education

Massachusetts Institute of Technology (MIT)

Doctor of Philosophy in Electrical Engineering and Computer Science

Advisor Prof. Dylan Hadfield-Menell

Cambridge, MA

Incoming, February 2027

Carnegie Mellon University (CMU)

Master of Science in Machine Learning

GPA 4.24/4.0

Advisors Prof. Virginia Smith, Prof. Aditi Raghunathan

Committee Service MSML Admissions Committee, MSML Social Committee, SCS Master's Advisory Committee

Pittsburgh, PA

December 2025

Worcester Polytechnic Institute (WPI)

Bachelor of Science in Computer Science, **2nd Major** in Mathematics, **Minor** in Robotics Engineering

GPA 4.0/4.0

Awards and Activities Graduation with High Distinction, Men's Varsity Rowing

Worcester, MA

May 2024

Selected Publications

AI Safety and Oversight

[1] **Kale, N.**, Zhang, C. B. C., Zhu, K., Aich, A., Rodriguez, P., Knight, C. Q., & Wang, Z. (2026). Reliable Weak-to-Strong Monitoring of LLM Agents. In *International Conference on Learning Representations (ICLR)*. **Oral**.

[2] **Kale, N.***, Portnoff, R*, Thaker, P., Simpson, M., Wang, R., Kuo, K., Yadav, C., & Smith, V. (2026). Position: Preventing AI-Generated CSAM Necessitates New Approaches to AI Safety. In *International Conference on Machine Learning (ICML)*. **Spotlight**.

[3] Wu, C. H.*, **Kale, N.***, & Raghunathan, A. (2025). Mitigating modal imbalance in multimodal reasoning. In *Conference on Language Modeling (CoLM)*.

Privacy and Auditing

[4] Thaker, P*, **Kale, N.***, Wu, Z. S., & Smith, V. (2026). Membership Inference Attacks for Unseen Classes. In *International Conference on Machine Learning (ICML)*.

[5] Thaker, P., Hu, S., **Kale, N.**, Maurya, Y., Wu, Z. S., & Smith, V. (2025). Position: LLM unlearning benchmarks are weak measures of progress. In *IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)* (pp. 520-533).

Industry Experience

Trajectory

Founding Member of Technical Staff

Second employee at Trajectory, developing algorithms and infrastructure for continual learning.

February 2026 – Present

Scale AI (Hosted by Christina Knight)

Security and Policy Research Lab Fellow

Engaged with global policymakers to solve real-world problems in agent robustness, AI control, and frontier risk evaluation.

October 2025 - January 2026

Scale AI (Hosted by Sail Wang & Julian Michael)

SEAL Research Scientist Intern

Developed novel chain-of-thought monitoring techniques to oversee superhuman agents with weaker LLM monitors; first-author paper accepted as an oral at ICLR 2026.

May 2025 - August 2025

GoDaddy Inc (Hosted by Harsh Nilesh Pathak)

Machine Learning Scientist Intern

Built GoDaddy's first agentic pipeline for CRM and prototyped an early agentic chatbot.

May 2024 - August 2024

Teaching

Teaching Assistant for 10701 Introduction to Machine Learning (CMU), CS541 Deep Learning (WPI)